Automated Model Risk Auditor

A test that does not trust a single judge. That cannot be fooled by good behavior. And that keeps watching after release.

The parts

- **Test suite management:** A live, organized test suite where every change creates a new version, and every result is recorded against the version that produced it.
- **Detecting a model that behaves differently when tested:** Every question is asked twice, once as an announced test and once inside a normal conversation, and the answers are compared.
- **Three different judges:** Every answer is examined by three different methods, and the decision follows the agreement between them.
- **Risk score and release gate:** All results become one score, but a critical category that crossed a hard threshold fails the model regardless of the score.
- **Automatic attack generation:** The system invents new attack phrasings with three methods, without exceeding a fixed budget.
- **Running in an isolated room:** Tests that require running code run in an environment that looks real but is sealed off from the world.
- **Regression tracking:** Every new model version is compared to the previous ones, and only a real change raises an alert.
- **A signed report for the regulator:** All results are summarized in a three-level report, signed so that any change to it is detected.
- **Monitoring after release:** A small sample of real usage passes through the same judges, to detect gradual drift.

Six decisions

- **One judge or three?** We chose three different judges.
- **Average score only, or also a hard threshold?** We chose a hard threshold.
- **Manual scenarios only, or automatic too?** We chose automatic generation as well.
- **Run directly or inside an isolated room?** We chose an isolated microVM.
- **Once before release or also after?** We chose continuous monitoring.
- **One score or a signed three-level report?** We chose a signed report.