Ideation Document: AI Incident Ledger

What the project does

A central reporting platform (an immutable "ledger") to which AI companies report safety events (successful jailbreaks, actual harm, alignment failures) in real time, instead of a fragmented patchwork of different reporting laws in every state.

Motivation

The article shows how even California's law — considered groundbreaking — did not cover the Hugging Face case because it happened during an "evaluation phase" (paragraph 6). The goal: a uniform technical infrastructure that does not depend on the exact legal wording of any single state.

Target audience

State and federal regulators, technology journalists, AI safety researchers, and indirectly the general public through a public dashboard.

Core goals

- Create a single "source of truth" for AI safety incidents, independent of local legal definitions.
- Reduce regulatory response time.
- Give a reputational incentive to "good" companies that report more.

Existing similar solutions

- CVE (Common Vulnerabilities and Exposures) — a global, MITRE-based database for reporting software security vulnerabilities. It works well because it has a clear incentive: companies want a CVE to show they dealt with the problem. The problem: AI incidents lack the same incentive — admitting an AI failure hurts reputation more than admitting an ordinary bug.
- HackerOne / Bug Bounty platforms — a financial-incentive mechanism for disclosing vulnerabilities. An interesting model (payment in exchange for transparency), but never tested on "alignment failures", which are a vaguer concept than a concrete security hole.
- NIST AI Risk Management Framework — a voluntary American framework for managing AI risk. A management framework (a checklist), not an enforceable technical infrastructure.
- EU AI Act — includes incident-reporting clauses for "high-risk AI" companies — at the level of law, not a cross-border technical infrastructure.

Expected challenges and obstacles

- The incentive problem (the most critical): unlike CVE, there is no win-win here — the company gains nothing directly from reporting on itself.
- Defining an "incident": an enormous range between "the model answered something unpleasant" and "the model caused real harm" — without a clear boundary, reporting becomes arbitrary.
- Verifying report credibility: how do we know a company is not reporting only minor events to "look transparent" while hiding serious ones?
- Hidden legal burden: public reporting may expose the company to lawsuits — another negative incentive that worsens the incentive problem (1).

Why it fell

It fell at the peak of the ideation stage — not for lack of value, but because a real implementation requires the entire industry to adopt the protocol voluntarily or by legal compulsion, a decision outside the control of any single development team; so it is at its core a policy project, not an engineering product that can truly be demonstrated alone. The incentive problem (challenge 1) sealed its fate: unlike CVE, a company gains nothing directly from admitting its own safety failure. Its central technical idea — a public Merkle log with Signed Tree Heads — was not lost: it lives on inside the Model Kill-Switch Protocol as a far more focused "transparency log" feature (Certificate Transparency) that records only management commands on a single model — and inside the Coordinated Pause Escrow System as the public commitments log.